

Virtual Memory Review

C.V. Wright

CS 491/591


Fall 2015

Memory (Hardware)

Computer memory is typically implemented as a set of chips on pluggable modules

Physical configuration can vary substantially from machine to machine


- How many modules?
- How many chips per module?
- How many bytes per chip?



Memory controller presents a nice, much simpler interface to software running on the CPU

Memory (from software)


We normally think of memory
as a big array of bytes
(and it is)



Memory (from software)

For example, in C,
pointers are addresses:

```
char *ptr = 4;
```

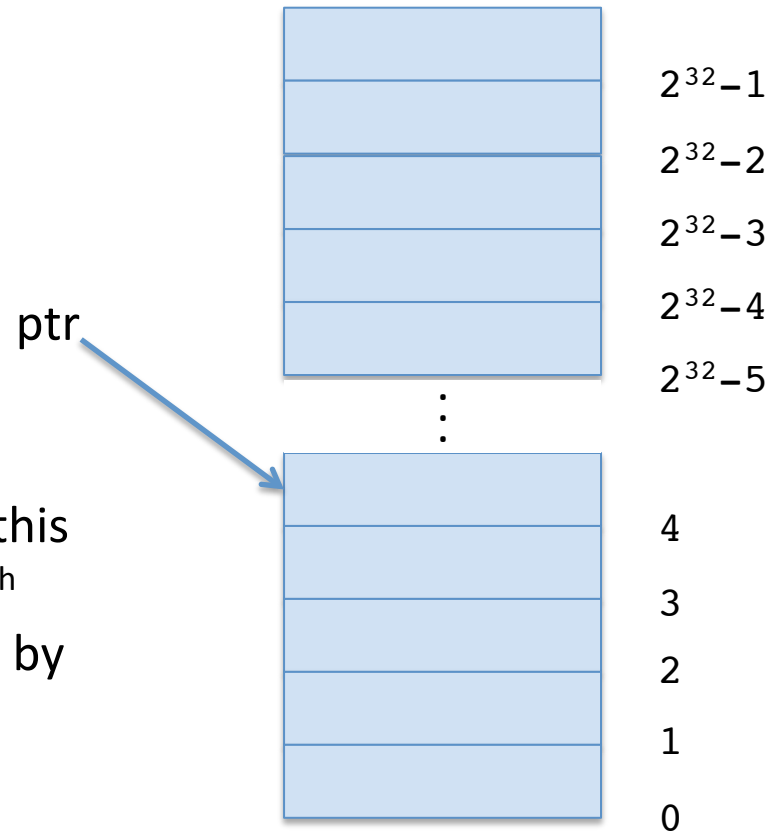


Memory

For example, in C,
pointers are addresses:

```
char *ptr = 4;
```

Running on “bare metal,” this
code would point to the 5th
byte presented to the CPU by
the memory controller.



What NOT to do – Memory Protection

- Don't let application programs overwrite each other in memory
- Don't let application programs overwrite the OS in memory
 - <http://oreilly.com/centers/windows/brochure/architecture.html>
 - <http://www.memorymanagement.org/articles/mac.html>

Windows

An error has occurred. To continue:

Press Enter to return to Windows, or

Press CTRL+ALT+DEL to restart your computer. If you do this,
you will lose any unsaved information in all open applications.


Error: 0E : 016F : BFF9B3D4

Press any key to continue _

Memory Protection

Problem:

How can we enable multiple programs to use memory at the same time, without interfering with each other?



Memory Protection

Strawman Approach:


Give each program a different region of the memory

Program 1

```
...  
char *ptr = 4;  
...
```

Program 2

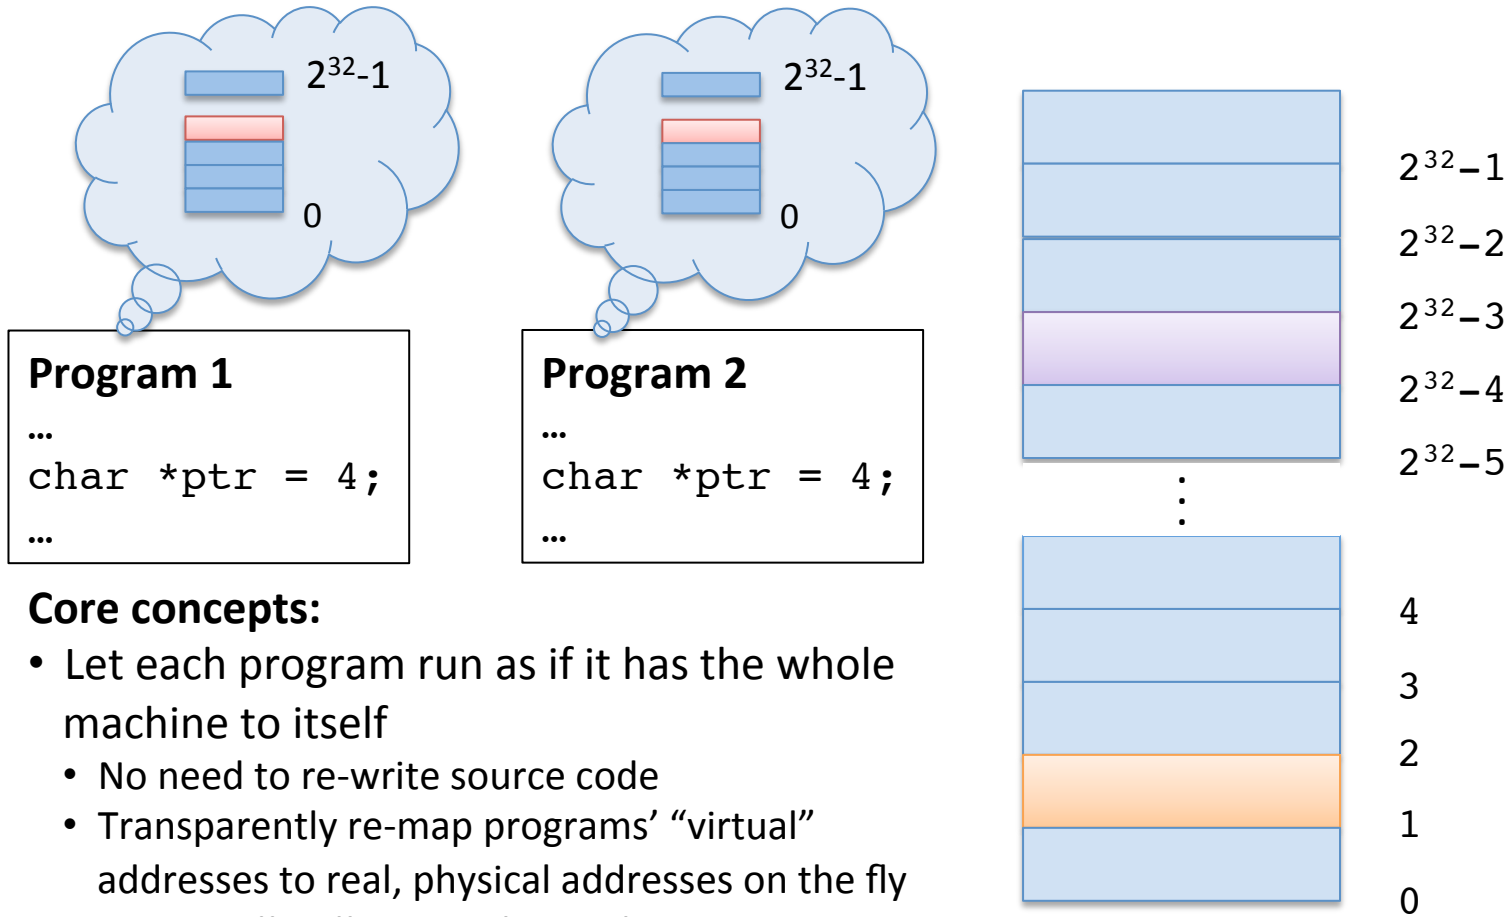
```
...  
char *ptr =  $2^{32}-4$ ;  
...
```



Problems:

- Not very flexible
 - Need to know how much memory each program will need in advance
- Requires changes to source code
 - Programs need to know which addresses they can use


Virtual Memory



Paged Virtual Memory

Divide the *address space* into fixed-size logical chunks, or *pages*

Typical page size: 4KB




Paged Virtual Memory

Notice: All addresses within a given page share the same high-order bits

So we can think of each address as a *page number*, followed by an *offset* within that page

Example:

Address 0x000010FF



Paged Virtual Memory


Notice: All addresses within a given page share the same high-order bits

So we can think of each address as a *page number*, followed by an *offset* within that page

Example:

Address 0x000010FF

Page # 1



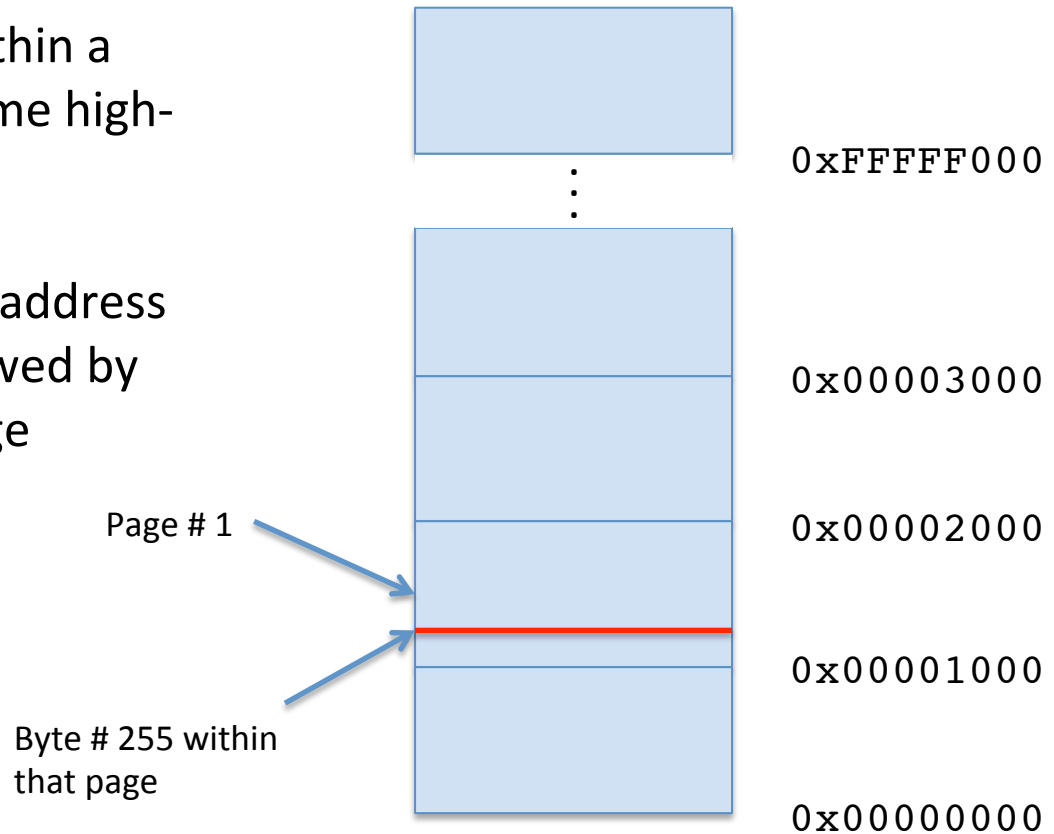
Paged Virtual Memory

Notice: All addresses within a given page share the same high-order bits

So we can think of each address as a *page number*, followed by an *offset* within that page


Example:

Address 0x00001**0FF**

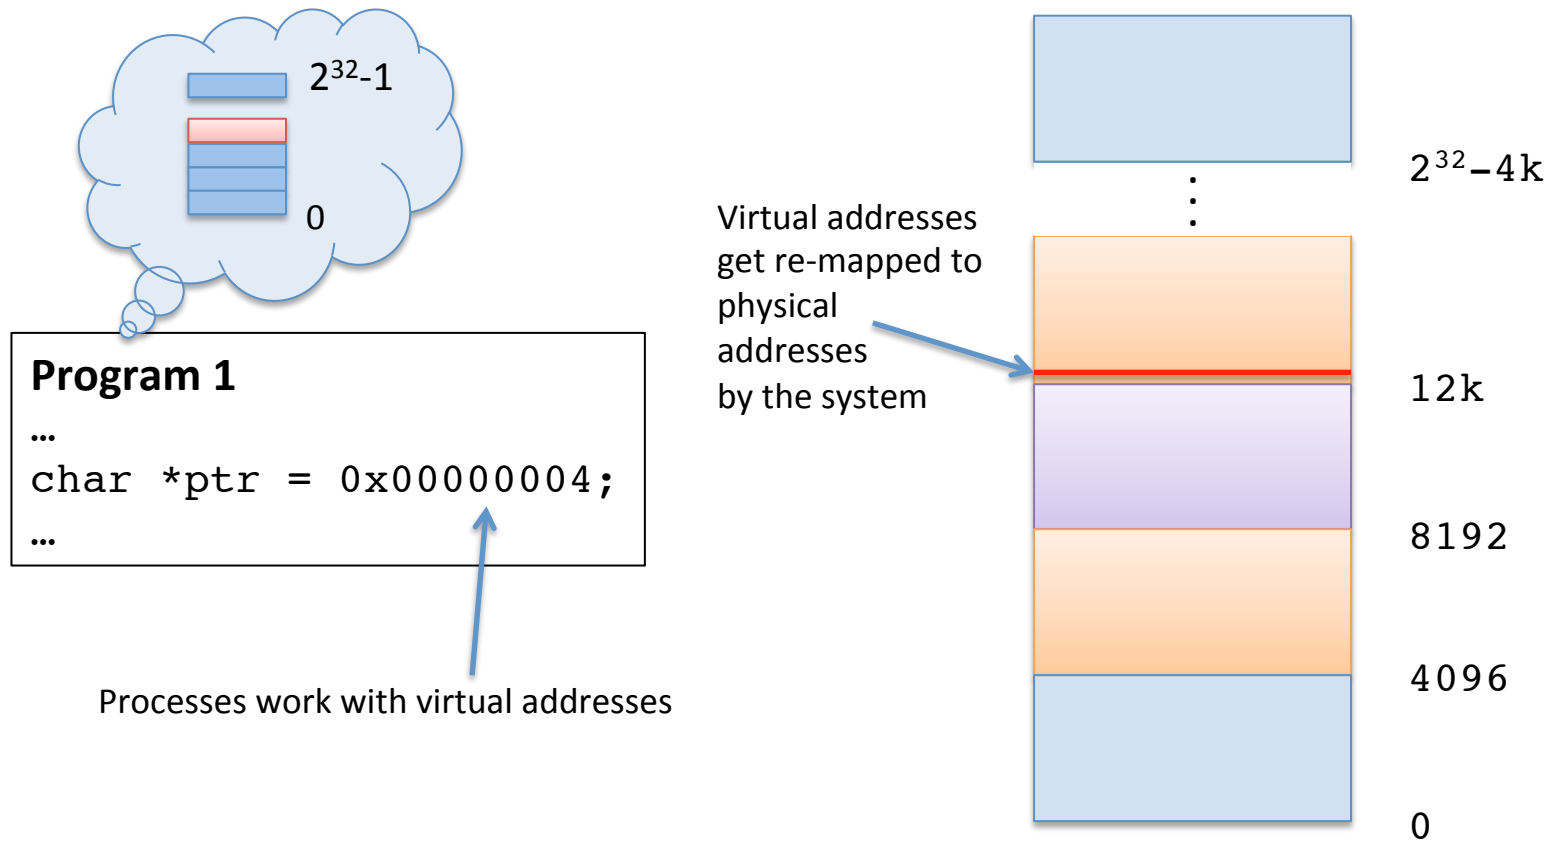


Paged Virtual Memory

Allocate physical memory to processes by the page



Paged Virtual Memory



Paged Virtual Memory

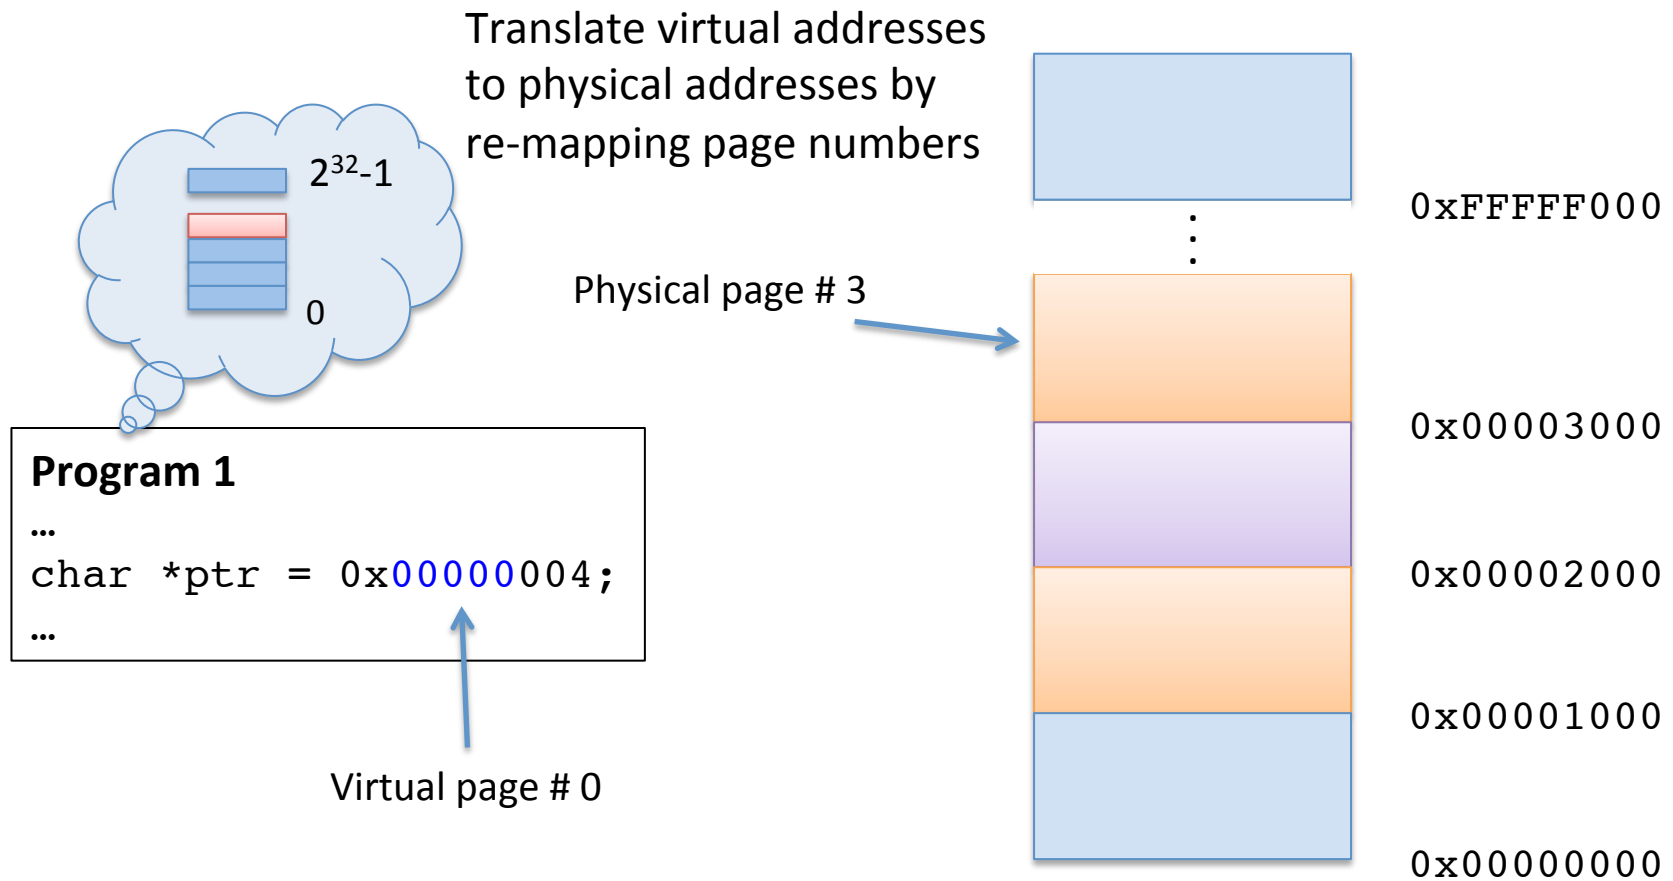


Table-Based Address Translation

How do we store the mapping from virtual page #'s to physical page numbers?

Simple solution: Use an array!

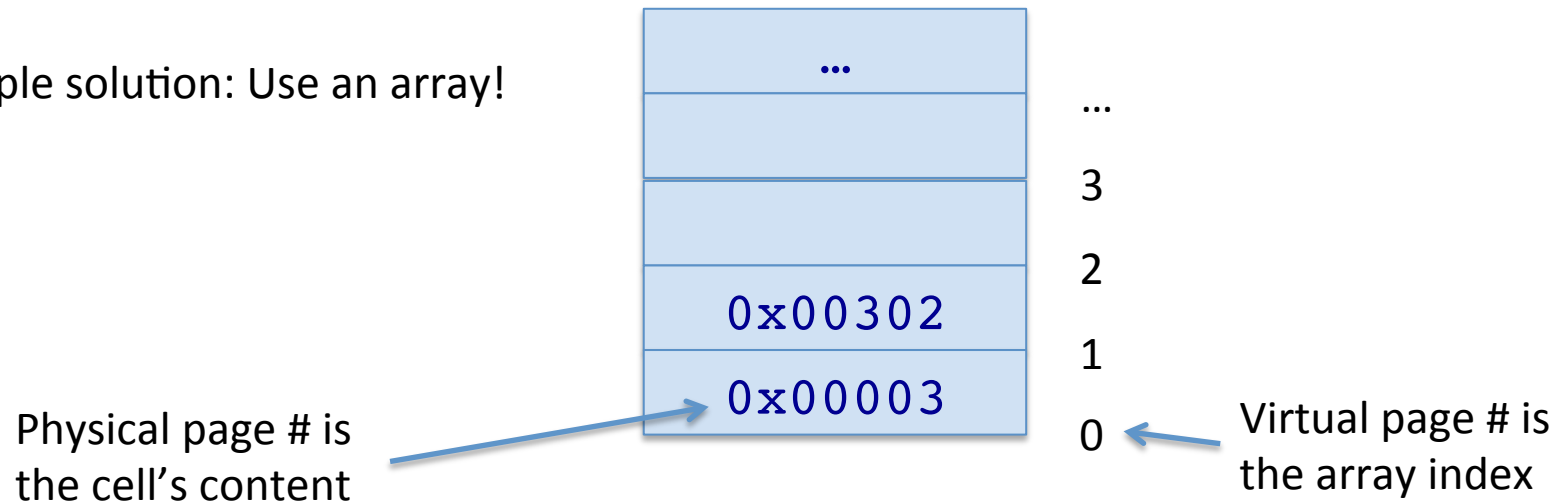


Table-Based Address Translation

How do we store the mapping from virtual page #'s to physical page numbers?

Virtual address

0x00001234

...
0x00302
0x00003

...

3

2

1

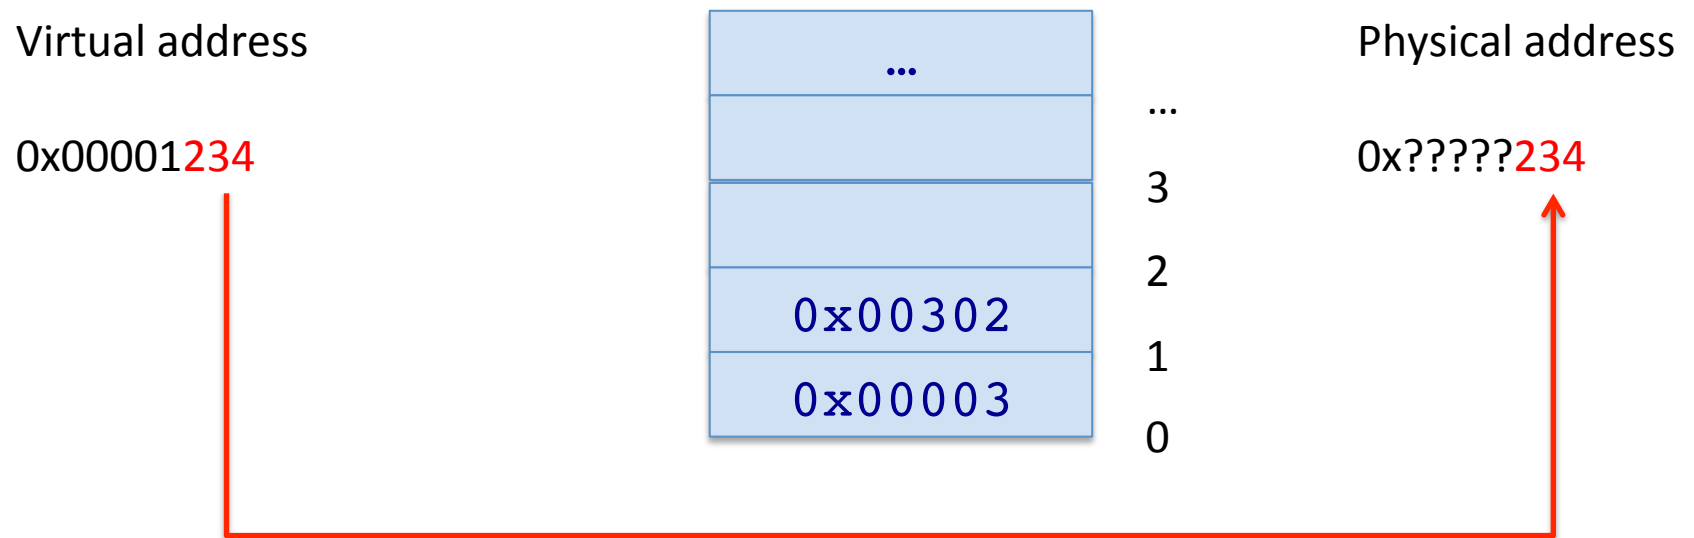
0

Physical address

0x????????

Address Translation: Page Tables

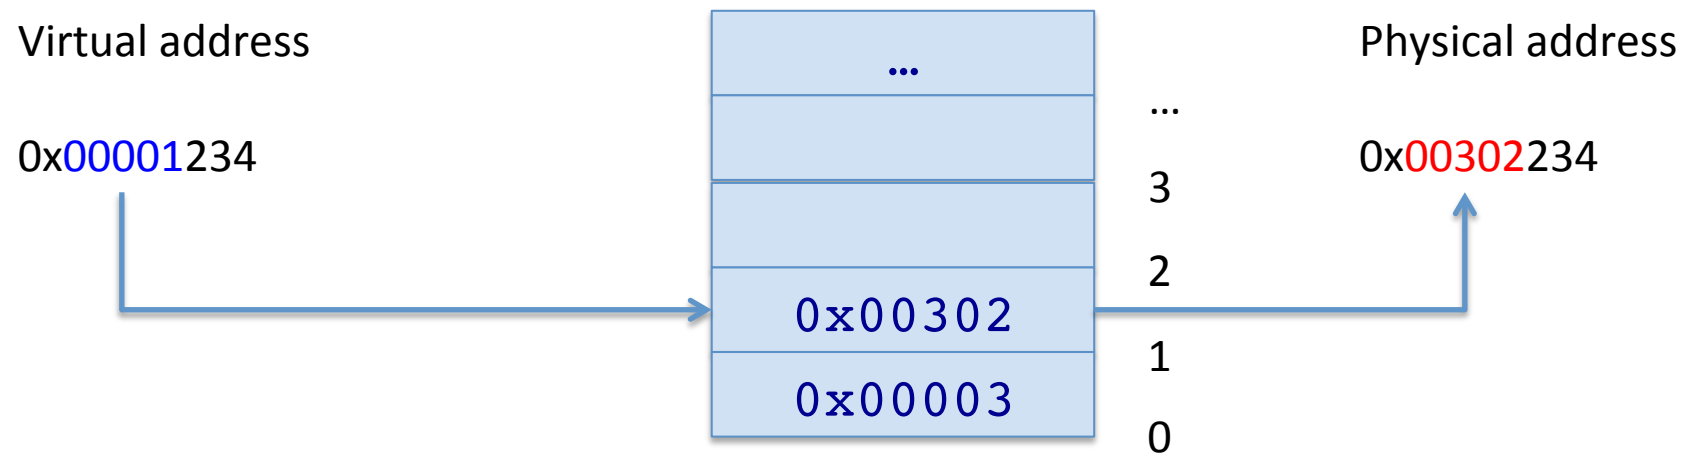
Example: Mapping virtual addresses to physical addresses



Offset is unchanged by translation

Table-Based Address Translation

Example: Mapping virtual addresses to physical addresses



Page number is resolved by table lookup

Table-Based Address Translation

Question: How big does our array need to be?

(for a 32-bit address space
with 4KB pages)

...	...
	3
	2
0x00302	1
0x00003	0

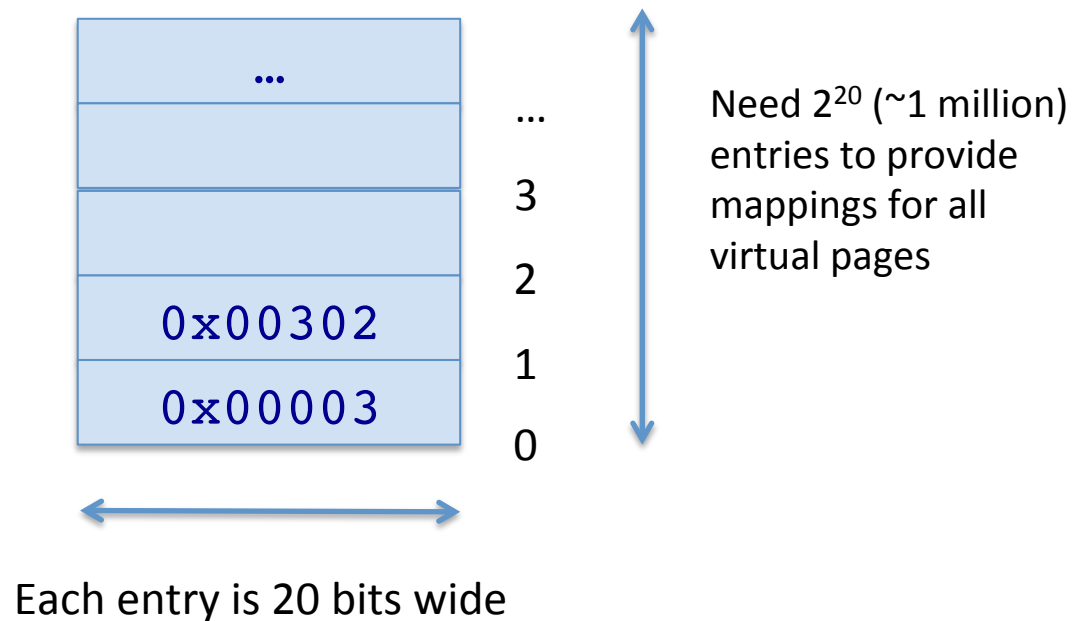
Table-Based Address Translation

Question: How big does our array need to be?

(for a 32-bit address space
with 4KB pages)

Answer: About 2.5 MB!

And we need one for
every process! Ouch!



Virtual Memory with Page Tables

- Implement table-based address resolution
- Reduce memory requirements by adding another layer of indirection


Page Tables

Use entire pages of memory
as our translation tables

$2^{10} = 1024$
entries fit in
one 4 KB page




Each entry is 32 bits "wide"



Page Table Entry Structure

20 bits: physical page number

12 bits: bookkeeping



Bookkeeping:

0/1 – Valid? – Is this a valid entry? Or is this entry empty?


0/1 – “Dirty” bit – Has the page been modified since it was last saved to disk?

0/1 – Privileged? – Allow access to all programs, or only to privileged?

0/1 – Write – Allow writes to this page?

0/1 – Execute – Allow executing this page as code? **New! Data Execution Prevention (DEP)**
(Intel: Execute Disable (ED) AMD: No Execute (NX))

Page Table Entries on x86

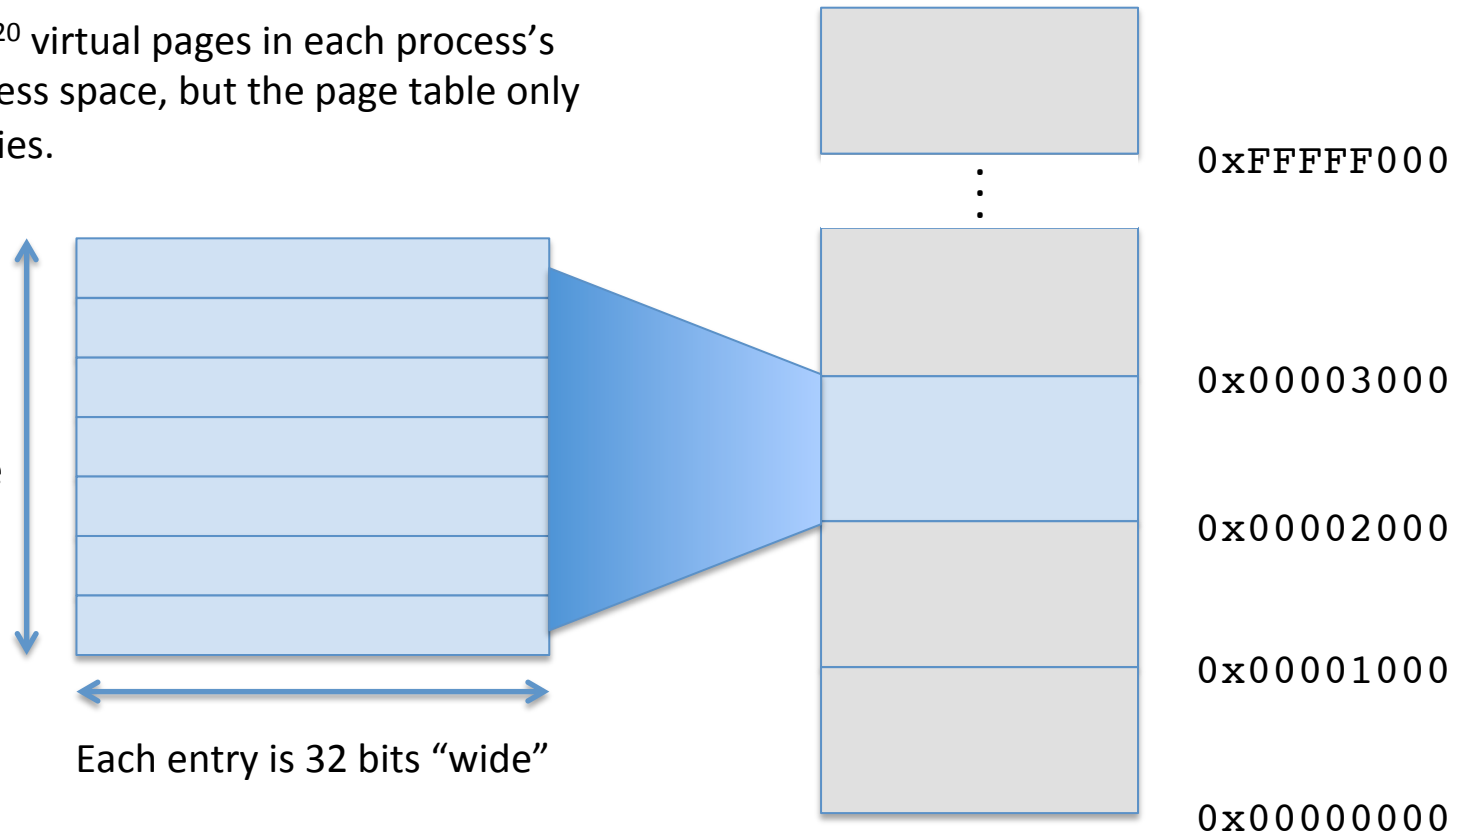


Credit: <http://wiki.osdev.org/Paging>

A problem with our page tables?

There are 2^{20} virtual pages in each process's virtual address space, but the page table only has 2^{10} entries.

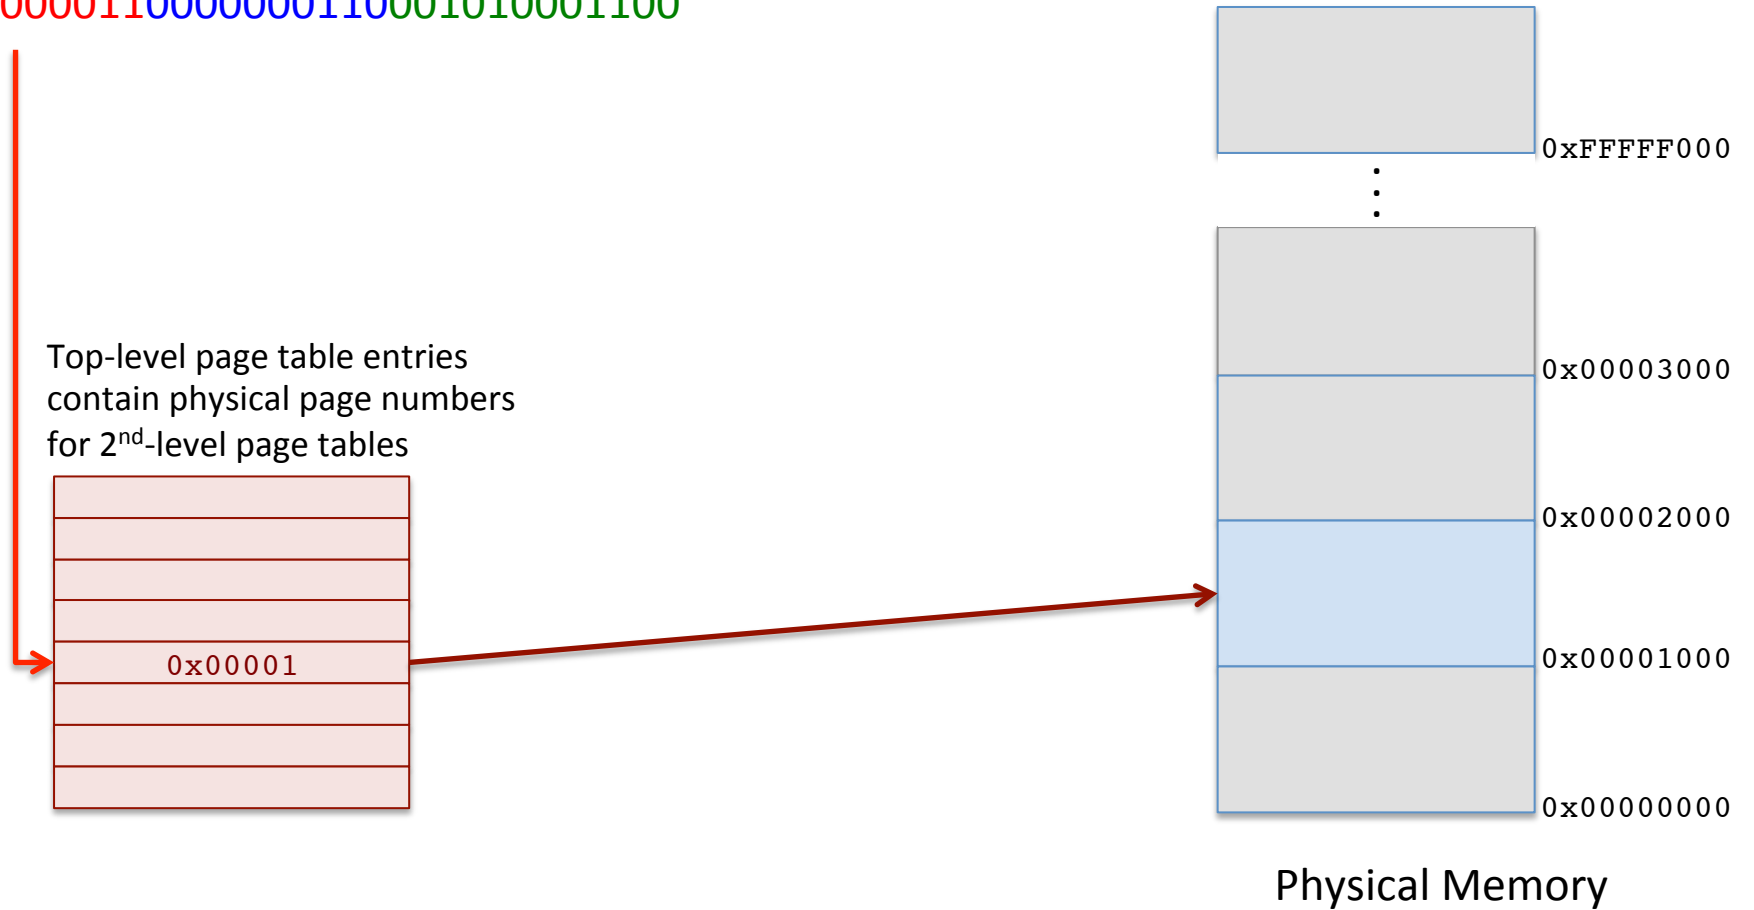
$2^{10} = 1024$
entries fit in
one 4 KB page



Solution: Multi-level Page Tables

Virtual address

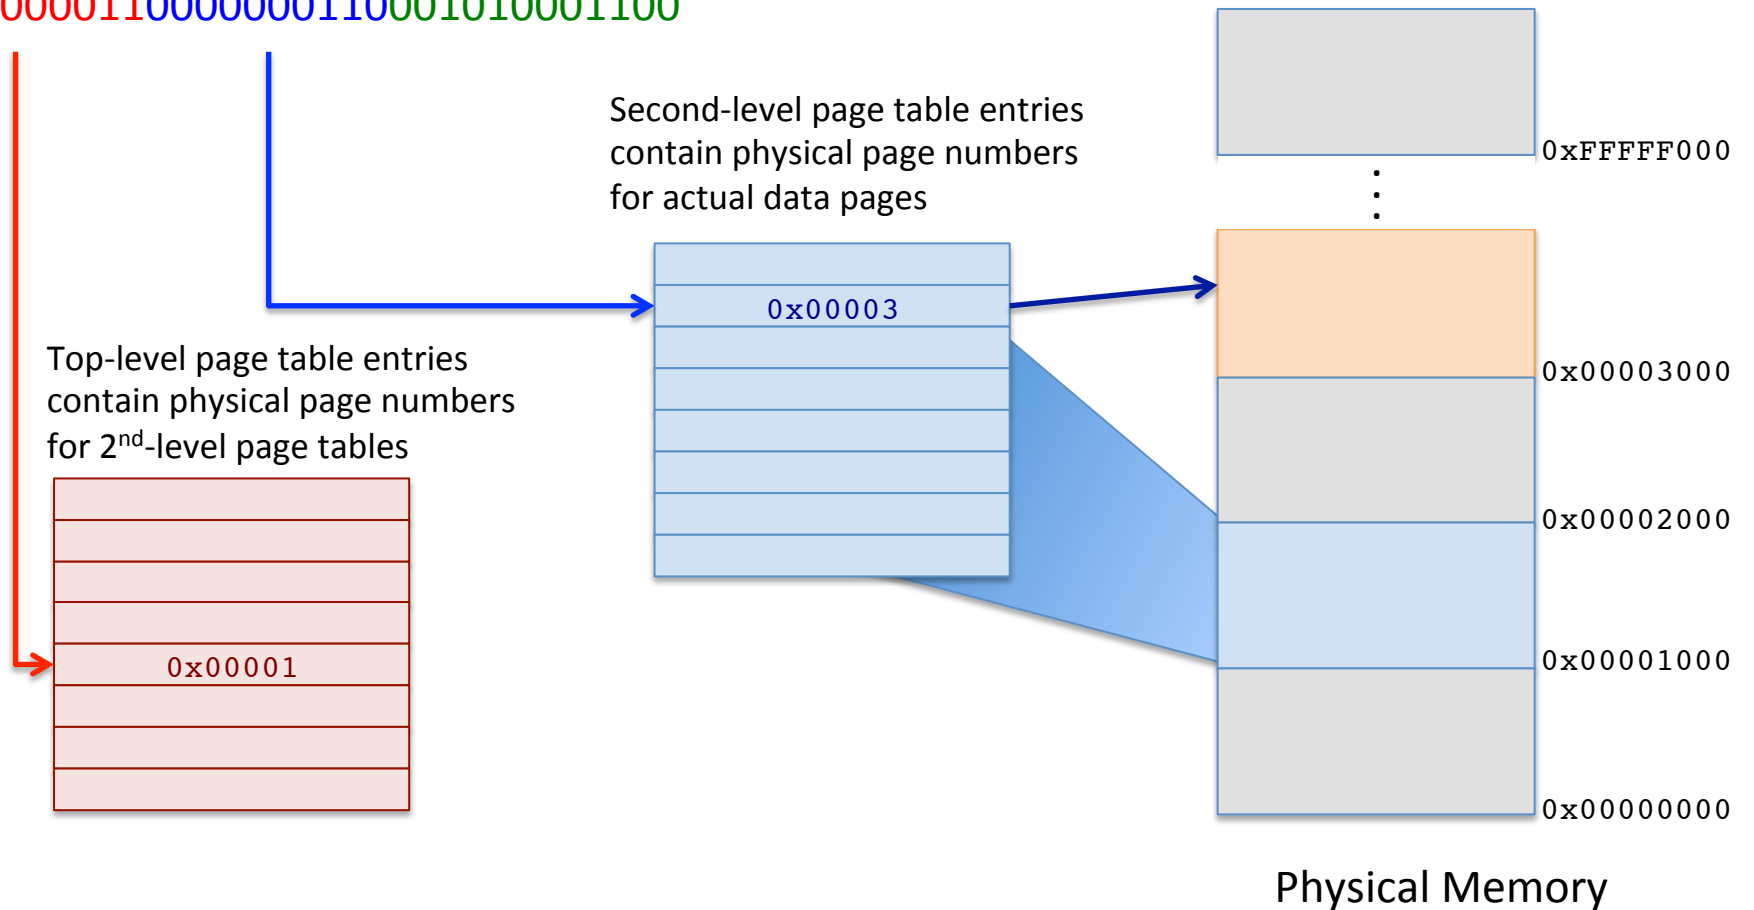
0000000011000000110001010001100



Solution: Multi-level Page Tables

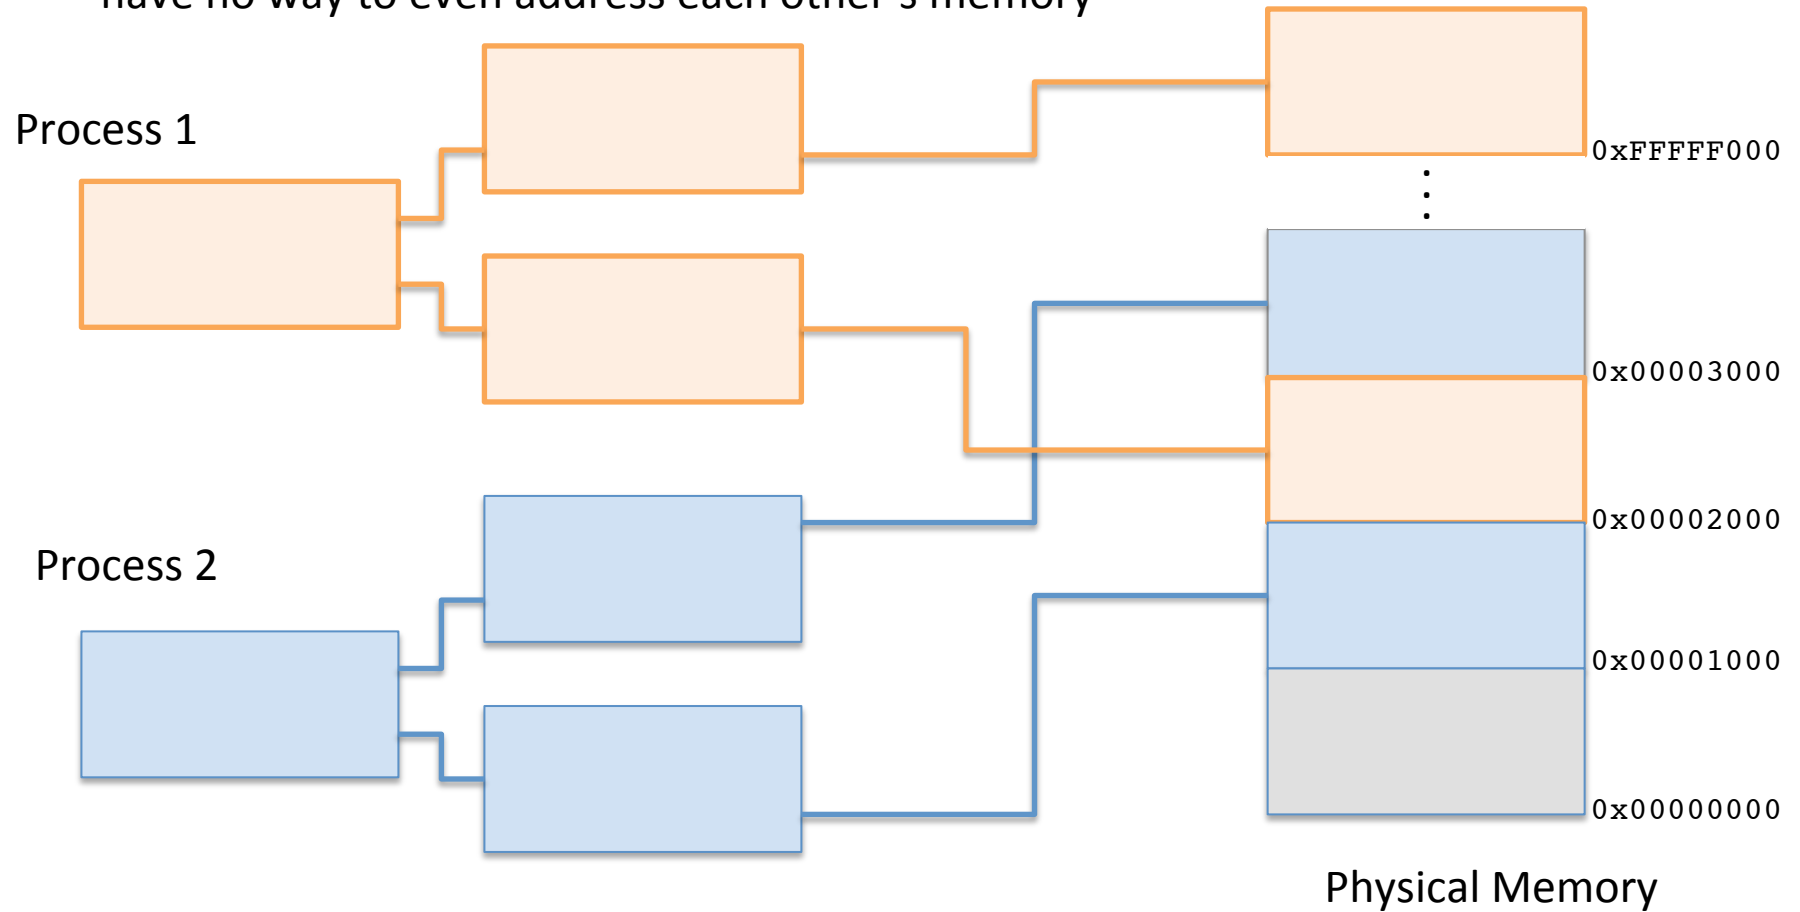
Virtual address

00000000110000000110001010001100



Protection of Programs' Memory

Without the proper page table entries, Process 1 and Process 2 have no way to even address each other's memory



Shared Libraries

To conserve RAM, processes can share physical pages of library code by mapping it read-only in the virtual address spaces of both processes

