

# Heap Overflows and Data Execution Prevention (DEP)

CS 491/591

Fall 2015

No more code on the stack!

# **DATA EXECUTION PREVENTION**

# Defense: What can we do?

1. Get programmers to write better code?
  - Good luck with that...
2. Get programmers to use a safe language?
  - Maybe someday, not today...
3. Modify the compiler?
4. Modify the OS and hardware?


# Data Execution Prevention (DEP) aka NX bit, aka W<sup>X</sup>

- Idea: Make data regions non-executable
- Old implementation: Use x86 segmentation
- New implementation: Use new hardware protection features

# Page Table Entry Structure

20 bits: physical page number

12 bits: bookkeeping



## Bookkeeping:

0/1 – Valid? – Is this a valid entry? Or is this entry empty?


0/1 – “Dirty” bit – Has the page been modified since it was last saved to disk?

0/1 – Privileged? – Allow access to all programs, or only to privileged?

0/1 – Write – Allow writes to this page?

0/1 – Execute – Allow executing this page as code? **New! Data Execution Prevention (DEP)**  
(Intel: Execute Disable (ED)      AMD: No Execute (NX) )

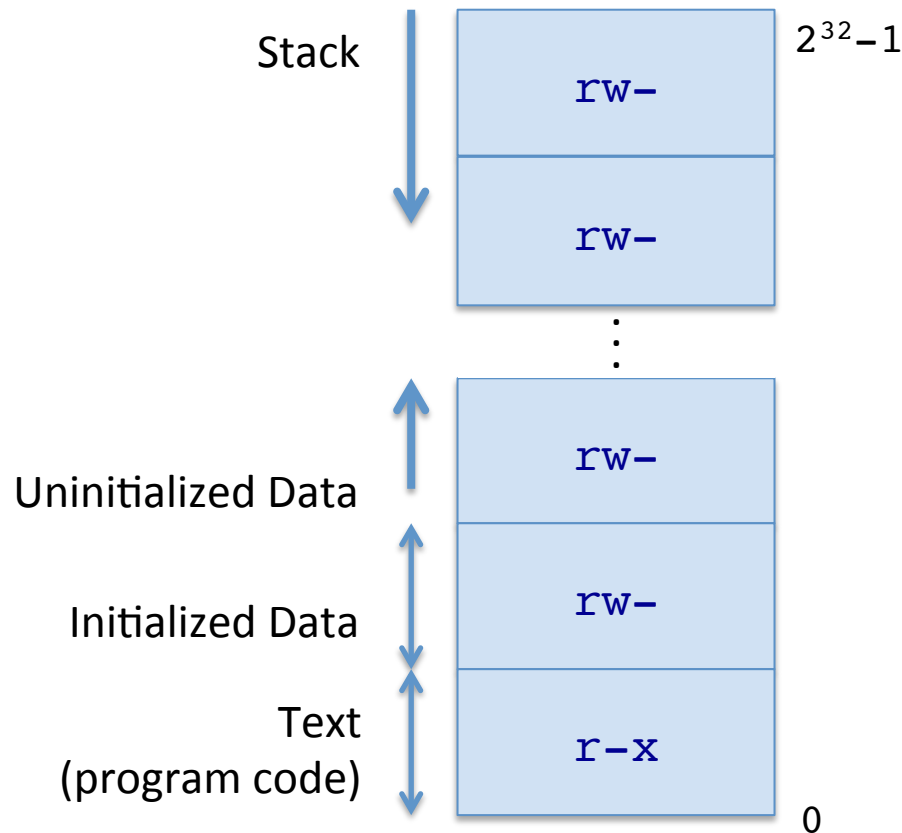
# Page Table Entries on x86



Credit: <http://wiki.osdev.org/Paging>

# Data Execution Prevention

## Process Virtual Address Space



No page has both write (w) and execute (x) permissions

Even if the attacker can inject code, he can't execute it!

# Remember Example 2?

```
#include <string.h>

void function(char *str) {
    char buffer[16];
    strcpy(buffer, str);
}


void main() {
    char large_string[256];
    int i;

    for(i=0; i < 255; i++)
        large_string[i]='A';


    function(large_string);
}
```




## Example 2 – After strcpy()




## Example 2 – leave (1)




## Example 2 – leave (2)



# Example 2 – ret




# Example 2 – ret



Hey! Wait just a minute!

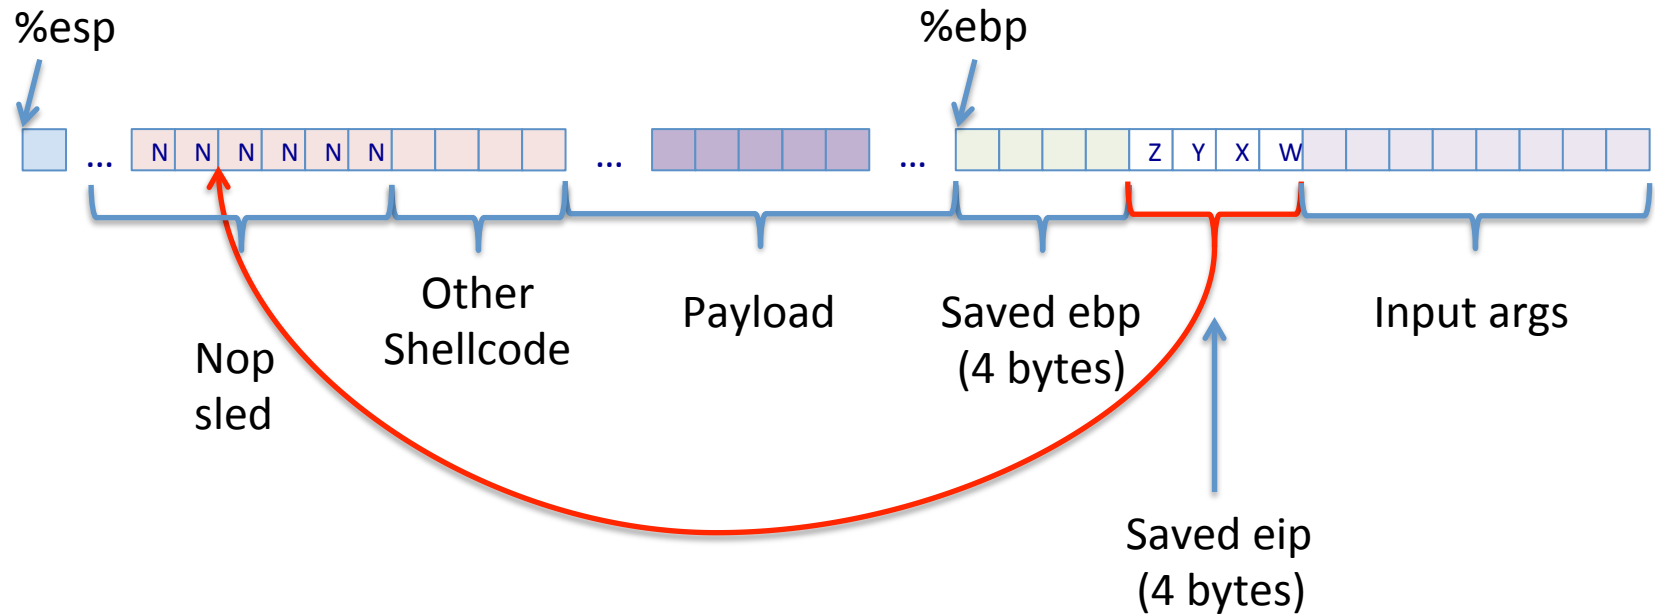
Do we have an entry  
in our page table for  
virtual page # 0x41414 ?

# Example 2 – ret

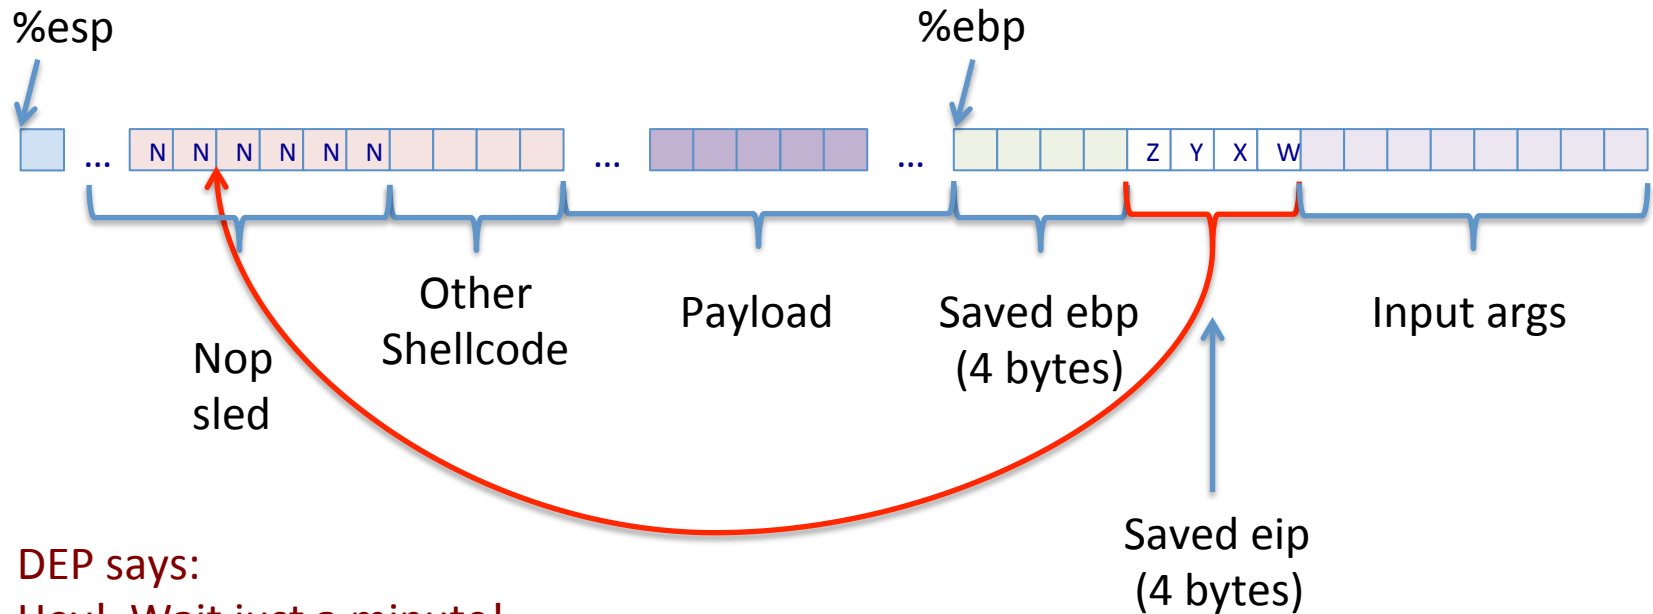


MMU says “No” → Segmentation Fault!

# Stack Buffer Overflow



# Stack Buffer Overflow with DEP



DEP says:  
Hey! Wait just a minute!

Do we have execute  
permissions in our page  
table for the virtual page  
with WXYZ? (the stack)

MMU says "No" → Segmentation Fault!



Smashing more than just the stack


# **HEAP OVERFLOWS**

# General Heap Overflow Problems

- Background
  - What is the heap?
  - Why is it a valuable target?
- Simple heap buffer overflow attacks
  - Changing filenames
  - Changing integers
  - Changing function pointers

# What is the heap?

# In-Memory Layout of a Process



# Example Program

```
#include <stdio.h>
#include <malloc.h>

int A;
int B;

int fcn(int depth) {
    return 0;
}

int main() {
    int x;
    char *buffer = (char *) malloc(128*sizeof(char));
    int *array = (int *) malloc(256*sizeof(int));

    fcn(10);

    return 0;
}
```

# Example Program

```
#include <stdio.h>
#include <malloc.h>
```

```
int A;
int B; } Global variables
```

Functions

```
int fcn(int depth) {
    return 0;
}

int main() {
    int x;
    char *buffer = (char *) malloc(128*sizeof(char));
    int *array = (int *) malloc(256*sizeof(int));

    fcn(10);

    return 0;
}
```

Local variable (data on the stack)

Dynamically-allocated variables  
(data on the heap)

Let's add some instrumentation to help us see what's going on

```
int main() {  
    int x;  
    char *buffer = (char *) malloc(128*sizeof(char));  
    int *array = (int *) malloc(256*sizeof(int));  
    void *main_ptr = main;  
    void *fcn_ptr = fcn;  
    void *x_ptr = &x;  
    void *printf_ptr = printf;  
    void *malloc_ptr = malloc;  
    void *A_ptr = &A;  
    void *B_ptr = &B;
```

Get addresses of  
variables in memory

```
    printf("Functions:\n");  
    printf("\t main() = %10p\n", main_ptr);  
    printf("\t fcn() = %10p\n", fcn_ptr);  
    printf("\t printf() = %10p\n", printf_ptr);  
    printf("\t malloc() = %10p\n", malloc_ptr);  
    printf("\n");  
    printf("Global Variables:\n");  
    printf("\t A = %10p\n", A_ptr);  
    printf("\t B = %10p\n", B_ptr);  
    printf("\n");  
    printf("Heap Variables:\n");  
    printf("\t buffer = %10p\n", buffer);  
    printf("\t array = %10p\n", array);  
    printf("\n");  
    printf("Stack Variables:\n");  
    printf("\t x = %10p\n", x_ptr);  
    printf("\n\n");
```

Print addresses  
in hex

```
    fcn(10);
```

```
    return 0;
```

```
}
```

```
[cvwright@ubuntu tmp]$ gcc -fno-stack-protector -o tracer2 tracer2.c  
[cvwright@ubuntu tmp]$  
[cvwright@ubuntu tmp]$ ./tracer2
```

Functions:

```
    main() = 0x804847e  
    fcn() = 0x8048474  
    printf() = 0x8048360  
    malloc() = 0x8048370
```

Global Variables:

```
    A = 0x804a02c  
    B = 0x804a028
```

Heap Variables:

```
    buffer = 0x804b008  
    array = 0x804b090
```

Stack Variables:

```
    x = 0xbffff6c8
```



```
[cvwright@ubuntu tmp]$ gcc -fno-stack-protector -o tracer2 tracer2.c
[cvwright@ubuntu tmp]$
[cvwright@ubuntu tmp]$ ./tracer2
```

Functions:

```
main() = 0x804847e
fcn() = 0x8048474
printf() = 0x8048360
malloc() = 0x8048370
```

Code at virtual page # 0x08048

Global Variables:

```
A = 0x804a02c
B = 0x804a028
```

Globals at virtual page # 0x0804a

Heap Variables:

```
buffer = 0x804b008
array = 0x804b090
```

Heap at virtual page # 0x0804b

Stack Variables:

```
x = 0xbffff6c8
```

Stack at virtual page # 0xbffff

# More instrumentation for function calls

```
int fcn(int arg) {  
    int rc;  
    char buf[5];  
    char *stuff = (char *) malloc(16*sizeof(char));  
  
    printf("depth = %2d    ", arg);  
    printf("arg = %10p    ", &arg);  
    printf("rc = %10p    ", &rc);  
    printf("buf = %10p    ", buf);  
    printf("stuff = %10p\n", stuff);  
  
    if(arg < 10)  
        rc = fcn(arg+1);  
    else  
        rc = 0;  
    free(stuff);  
    return rc;  
}
```

## Global Variables:

A = 0x804a034

B = 0x804a030

## Heap Variables:

buffer = 0x804b008

array = 0x804b090

## Stack Variables:

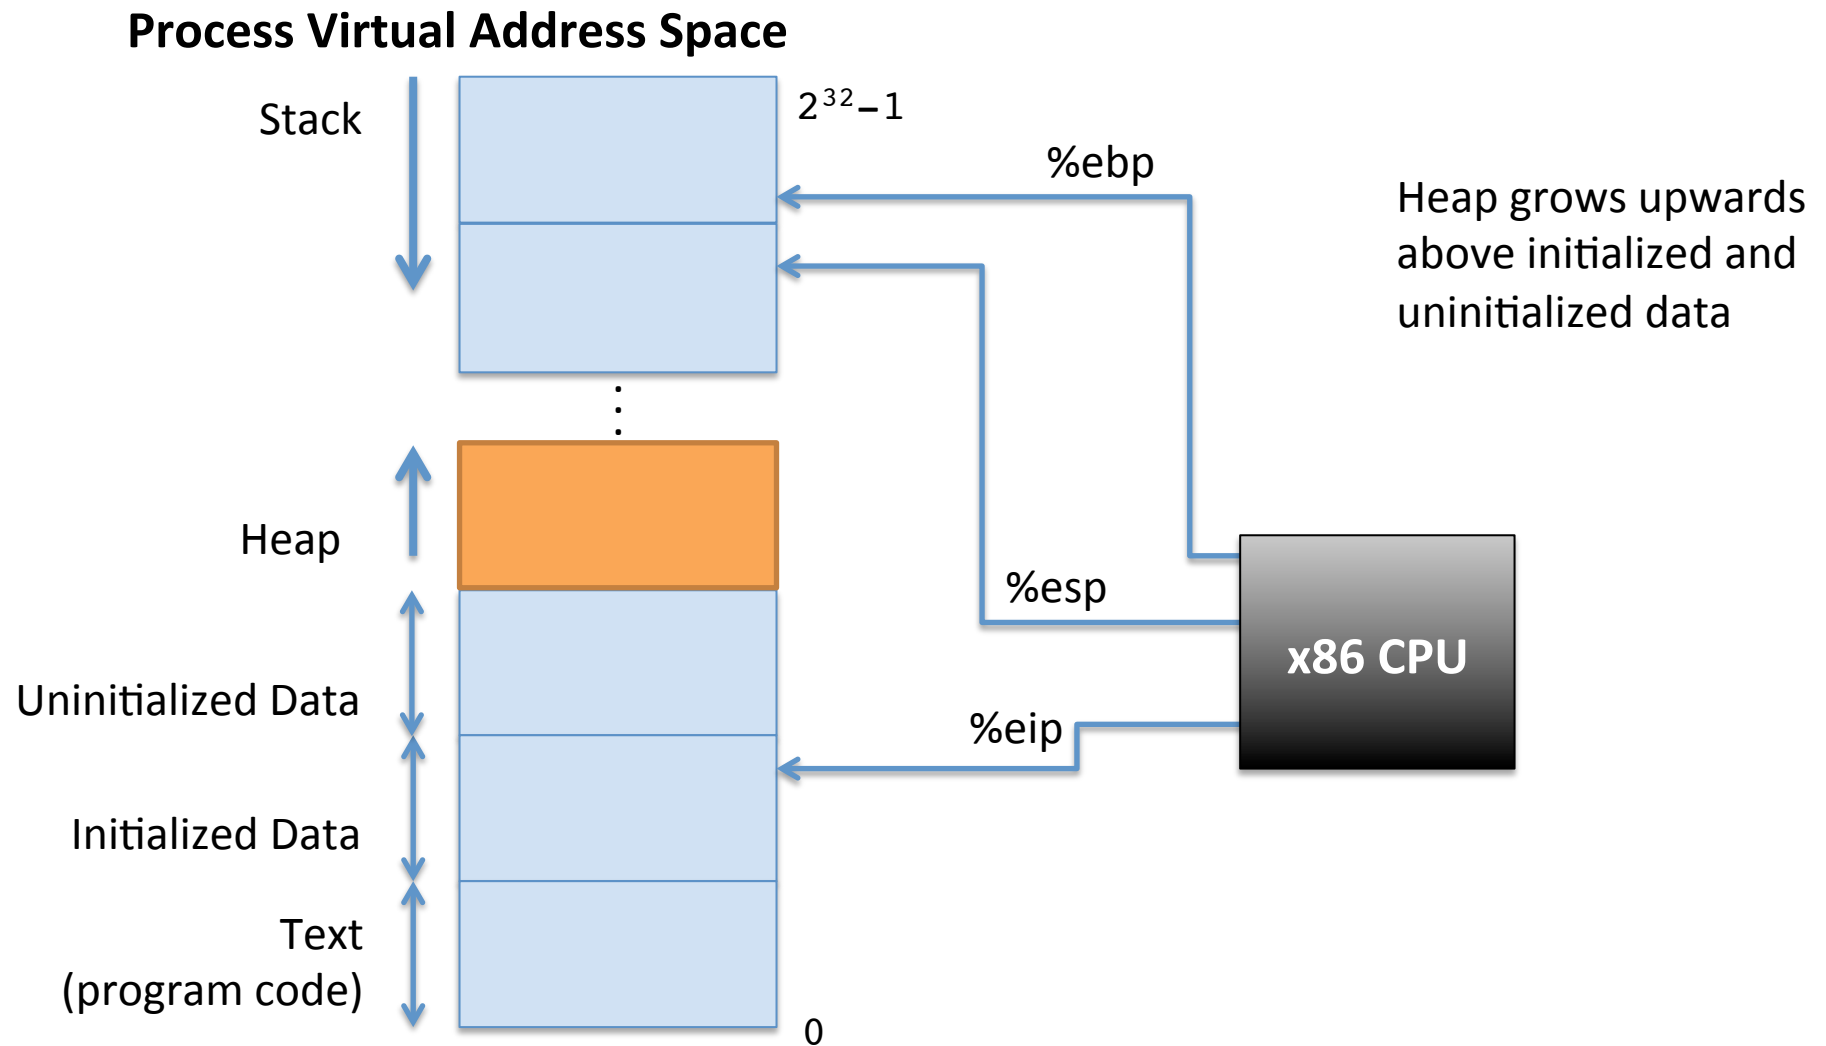
x = 0xbffff6c4

Stack grows  
"downwards"

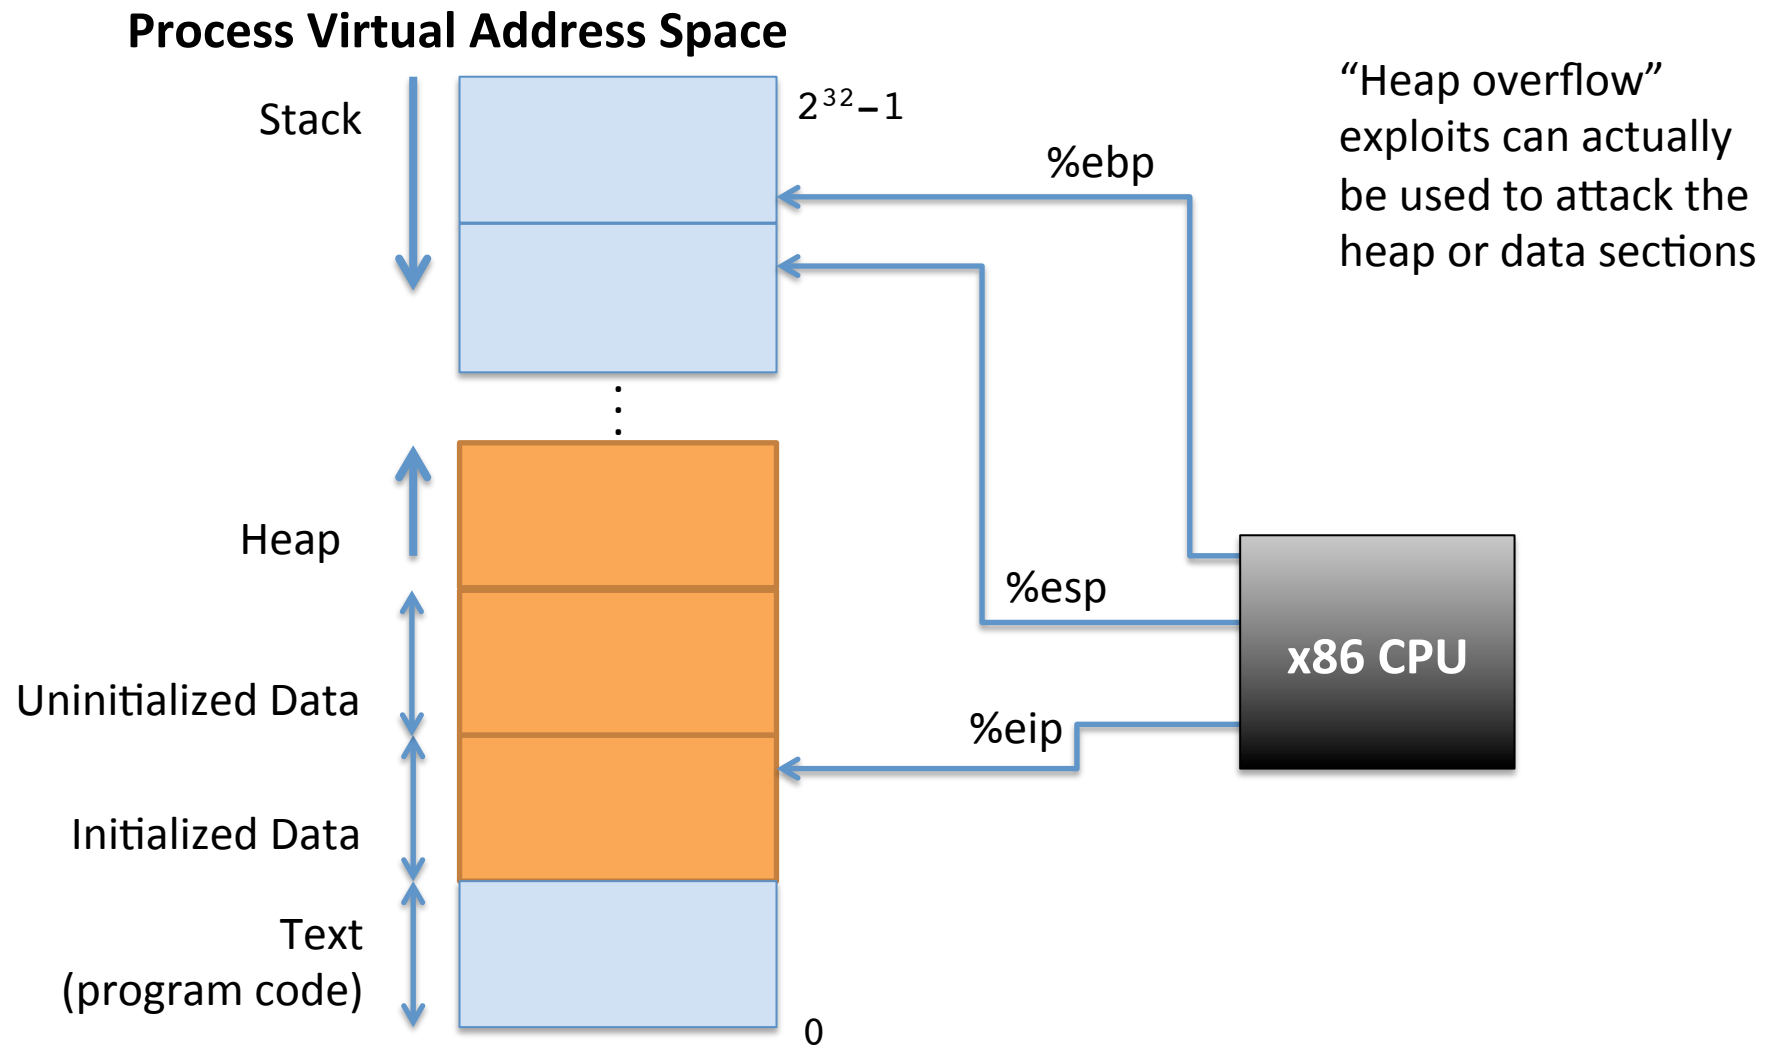
Heap grows  
"upwards"

depth = 0	arg = 0xbffff6b0	rc = 0xbffff698	buf = 0xbffff693	stuff = 0x804b498
depth = 1	arg = 0xbffff680	rc = 0xbffff668	buf = 0xbffff663	stuff = 0x804b4b0
depth = 2	arg = 0xbffff650	rc = 0xbffff638	buf = 0xbffff633	stuff = 0x804b4c8
depth = 3	arg = 0xbffff620	rc = 0xbffff608	buf = 0xbffff603	stuff = 0x804b4e0
depth = 4	arg = 0xbffff5f0	rc = 0xbffff5d8	buf = 0xbffff5d3	stuff = 0x804b4f8
depth = 5	arg = 0xbffff5c0	rc = 0xbffff5a8	buf = 0xbffff5a3	stuff = 0x804b510
depth = 6	arg = 0xbffff590	rc = 0xbffff578	buf = 0xbffff573	stuff = 0x804b528
depth = 7	arg = 0xbffff560	rc = 0xbffff548	buf = 0xbffff543	stuff = 0x804b540
depth = 8	arg = 0xbffff530	rc = 0xbffff518	buf = 0xbffff513	stuff = 0x804b558
depth = 9	arg = 0xbffff500	rc = 0xbffff4e8	buf = 0xbffff4e3	stuff = 0x804b570
depth = 10	arg = 0xbffff4d0	rc = 0xbffff4b8	buf = 0xbffff4b3	stuff = 0x804b588

# In-Memory Layout of a Process



# In-Memory Layout of a Process



# Why is the heap a valuable target?

Interviewer: *Well, can you... blow up the world?*

The Tick: *Egad! I hope not! That's where I keep all my stuff!*



From *The Tick* animated series (1994)

# Why is the heap a valuable target?

- That's where our programs keep (most of) their stuff
  - Global variables (really in the data sections)
  - Large, global data structures
  - Any data with dynamic size

# Simple Heap-Based Overflows

- Insecure use of the heap allows modification of other heap data
  - Overflow into a string → Change filename
  - Overflow into an int → Change user id
  - Overflow into function pointer → Change program's control flow



### Global Variables:

A = 0x804a034  
B = 0x804a030

### Heap Variables:

buffer = 0x804b008  
array = 0x804b090

### Stack Variables:

x = 0xbffff6c4

Stack grows  
"downwards"

Heap grows  
"upwards"

depth = 0	arg = 0xbffff6b0	rc = 0xbffff698	buf = 0xbffff693	stuff = 0x804b498
depth = 1	arg = 0xbffff680	rc = 0xbffff668	buf = 0xbffff663	stuff = 0x804b4b0
depth = 2	arg = 0xbffff650	rc = 0xbffff638	buf = 0xbffff633	stuff = 0x804b4c8
depth = 3	arg = 0xbffff620	rc = 0xbffff608	buf = 0xbffff603	stuff = 0x804b4e0
depth = 4	arg = 0xbffff5f0	rc = 0xbffff5d8	buf = 0xbffff5d3	stuff = 0x804b4f8
depth = 5	arg = 0xbffff5c0	rc = 0xbffff5a8	buf = 0xbffff5a3	stuff = 0x804b510
depth = 6	arg = 0xbffff590	rc = 0xbffff578	buf = 0xbffff573	stuff = 0x804b528
depth = 7	arg = 0xbffff560	rc = 0xbffff548	buf = 0xbffff543	stuff = 0x804b540
depth = 8	arg = 0xbffff530	rc = 0xbffff518	buf = 0xbffff513	stuff = 0x804b558
depth = 9	arg = 0xbffff500	rc = 0xbffff4e8	buf = 0xbffff4e3	stuff = 0x804b570
depth = 10	arg = 0xbffff4d0	rc = 0xbffff4b8	buf = 0xbffff4b3	stuff = 0x804b588

Why does the heap address increase by 0x18 ?

# More instrumentation for function calls

```
int fcn(int arg) {  
    int rc;  
    char buf[5];  
    char *stuff = (char *) malloc(16*sizeof(char));  
  
    printf("depth = %2d    ", arg);  
    printf("arg = %10p    ", &arg);  
    printf("rc = %10p    ", &rc);  
    printf("buf = %10p    ", buf);  
    printf("stuff = %10p\n", stuff);  
  
    if(arg < 10)  
        rc = fcn(arg+1);  
    else  
        rc = 0;  
    free(stuff);  
    return rc;  
}
```

After all, we only asked  
for 16 (0x10) bytes

# Example Heap Overflow

```
#include <stdio.h>
#include <malloc.h>

int main(int argc, char* argv[])
{
    char *username = (char *) malloc(16*sizeof(char));
    char *filename = (char *) malloc(16*sizeof(char));

    sprintf(filename, "%s.txt", argv[1]);
    if(!strcmp(filename, "flag.txt")){
        puts("FLAG ACCESS DENIED");
        return 1;
    }

    printf("Enter username for access to file [%s]\n", filename);
    scanf("%s", username);
```

# Attacking malloc() and free()

- How does malloc() work?
  - Chunks
  - Doubly-linked lists
- Abusing free() – the unlink() attack

# Malloc

- Malloc – An interface between the low-level POSIX OS interface and high-level C code
- POSIX – brk and sbrk system calls
  - Ask the OS to map more virtual memory at the top of the heap
  - Page-level granularity (e.g. 4KB at a time)
  - Managing already-allocated memory is up to the application itself

# Malloc

- Malloc's (and free's) job
  - Keep track of heap memory provided by the OS
  - Ask for more when we run out of space
  - Find the best location to satisfy each request from the application code

# Malloc

- How does malloc do it?
  - Sort of like a slab allocator from OS class (CS 333)
  - Tracks free regions using linked lists
  - Stores linked list pointers intermixed with data allocated for the application

# Malloc Implementation

```
#define INTERNAL_SIZE_T size_t

struct malloc_chunk {
    INTERNAL_SIZE_T prev_size;
    INTERNAL_SIZE_T size;
    struct malloc_chunk * fd;
    struct malloc_chunk * bk;
};
```



# Malloc Implementation

- An allocated chunk looks like this:

[illegible]

# Malloc Implementation

- Free chunks are stored in circular doubly-linked lists (described in 3.4.2) and look like this:

[illegible]

### Global Variables:

A = 0x804a034  
B = 0x804a030

### Heap Variables:

buffer = 0x804b008  
array = 0x804b090

### Stack Variables:

x = 0xbffff6c4

Stack grows  
"downwards"

Heap grows  
"upwards"

depth = 0	arg = 0xbffff6b0	rc = 0xbffff698	buf = 0xbffff693	stuff = 0x804b498
depth = 1	arg = 0xbffff680	rc = 0xbffff668	buf = 0xbffff663	stuff = 0x804b4b0
depth = 2	arg = 0xbffff650	rc = 0xbffff638	buf = 0xbffff633	stuff = 0x804b4c8
depth = 3	arg = 0xbffff620	rc = 0xbffff608	buf = 0xbffff603	stuff = 0x804b4e0
depth = 4	arg = 0xbffff5f0	rc = 0xbffff5d8	buf = 0xbffff5d3	stuff = 0x804b4f8
depth = 5	arg = 0xbffff5c0	rc = 0xbffff5a8	buf = 0xbffff5a3	stuff = 0x804b510
depth = 6	arg = 0xbffff590	rc = 0xbffff578	buf = 0xbffff573	stuff = 0x804b528
depth = 7	arg = 0xbffff560	rc = 0xbffff548	buf = 0xbffff543	stuff = 0x804b540
depth = 8	arg = 0xbffff530	rc = 0xbffff518	buf = 0xbffff513	stuff = 0x804b558
depth = 9	arg = 0xbffff500	rc = 0xbffff4e8	buf = 0xbffff4e3	stuff = 0x804b570
depth = 10	arg = 0xbffff4d0	rc = 0xbffff4b8	buf = 0xbffff4b3	stuff = 0x804b588

### Global Variables:

A = 0x804a034  
B = 0x804a030

### Heap Variables:

buffer = 0x804b008  
array = 0x804b090

### Stack Variables:

x = 0xbffff6c4

Stack grows  
"downwards"

Heap grows  
"upwards"

depth = 0	arg = 0xbffff6b0	rc = 0xbffff698	buf = 0xbffff693	stuff = 0x804b498
depth = 1	arg = 0xbffff680	rc = 0xbffff668	buf = 0xbffff663	stuff = 0x804b4b0
depth = 2	arg = 0xbffff650	rc = 0xbffff638	buf = 0xbffff633	stuff = 0x804b4c8
depth = 3	arg = 0xbffff620	rc = 0xbffff608	buf = 0xbffff603	stuff = 0x804b4e0
depth = 4	arg = 0xbffff5f0	rc = 0xbffff5d8	buf = 0xbffff5d3	stuff = 0x804b4f8
depth = 5	arg = 0xbffff5c0	rc = 0xbffff5a8	buf = 0xbffff5a3	stuff = 0x804b510
depth = 6	arg = 0xbffff590	rc = 0xbffff578	buf = 0xbffff573	stuff = 0x804b528
depth = 7	arg = 0xbffff560	rc = 0xbffff548	buf = 0xbffff543	stuff = 0x804b540
depth = 8	arg = 0xbffff530	rc = 0xbffff518	buf = 0xbffff513	stuff = 0x804b558
depth = 9	arg = 0xbffff500	rc = 0xbffff4e8	buf = 0xbffff4e3	stuff = 0x804b570
depth = 10	arg = 0xbffff4d0	rc = 0xbffff4b8	buf = 0xbffff4b3	stuff = 0x804b588

Why does the heap address increase by 0x18 ?

# More instrumentation for function calls

```
int fcn(int arg) {  
    int rc;  
    char buf[5];  
    char *stuff = (char *) malloc(16*sizeof(char));  
  
    printf("depth = %2d    ", arg);  
    printf("arg = %10p    ", &arg);  
    printf("rc = %10p    ", &rc);  
    printf("buf = %10p    ", buf);  
    printf("stuff = %10p\n", stuff);  
  
    if(arg < 10)  
        rc = fcn(arg+1);  
    else  
        rc = 0;  
    free(stuff);  
    return rc;  
}
```

After all, we only asked  
for 16 (0x10) bytes

# More instrumentation for function calls

```
int fcn(int arg) {  
    int rc;  
    char buf[5];  
    char *stuff = (char *) malloc(16*sizeof(char));  
  
    printf("depth = %2d    ", arg);  
    printf("arg = %10p    ", &arg);  
    printf("rc = %10p    ", &rc);  
    printf("buf = %10p    ", buf);  
    printf("stuff = %10p\n", stuff);  
  
    if(arg < 10)  
        rc = fcn(arg+1);  
    else  
        rc = 0;  
    free(stuff);  
    return rc;  
}
```

After all, we only asked  
for 16 (0x10) bytes

Answer:  
Malloc needs at least 8 bytes for  
its own internal book keeping!

# Heap Overflows and Malloc

- Free chunks are stored in circular doubly linked lists (described in 3.4.2) and look like this:

[illegible]

# The unlink attack

```
#define unlink( P, BK, FD ) {
    BK = P->bk;
    FD = P->fd;
    FD->bk = BK;
    BK->fd = FD;
}
```

///

**Basic idea: Overflow the preceding block to overwrite this block's bk and fd fields.**



# The unlink attack

```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

Malloc interprets the 4 bytes in fd as an address.

For example, this could be chosen so it's (close to) a function pointer or a saved %ebp.

# The unlink attack

```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

Then unlink writes to memory at a short offset from this address.

Here we cause it to actually overwrite the function pointer or saved eip.

# The unlink attack

```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

The value that it writes is taken from the 4 bytes of bk.

For our attack, BK could be the address of our shellcode.

# The unlink attack

```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

One slight complication. 4 bytes of the region at BK get clobbered!

## Now we have garbage in our shellcode. What to do?

# The unlink attack

```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```

\  
\  
\  
\  
\

One slight complication. 4 bytes of the region at BK get clobbered!

Now we have garbage in our shellcode. What to do?

One solution: Shellcode must contain a jmp to “jump over” the 4 bytes that get clobbered.